

# FDPS: Framework for Developing Particle Simulator

Ataru Tanikawa

(University of Tokyo / RIKEN Advanced Institute for  
Computational Science)

Collaborators

Masaki Iwasawa, Natsuki Hosono, Keigo Nitadori, Takayuki  
Muranushi, Junichiro Makino

(RIKEN Advanced Institute for Computational Science)

Oct. 30 2015 天体形成研究会 @ 筑波大学CCS

# What is FDPS ?

- FDPS is Short for “Framework for Developing Particle Simulator”.
- Using FDPS makes it easier to develop massively parallel particle simulation codes.
- Particle simulations include gravitational N-body, SPH, molecular dynamics, granular dynamics, etc.
- Governing equation:

$$\frac{d\vec{u}_i}{dt} = \vec{g} \left( \sum_j^N \vec{f}(\vec{u}_i, \vec{u}_j), \vec{u}_i \right)$$


Quantity vector of a particle

Function transforming particle's quantity to its time derivatives

Pairwise interaction function

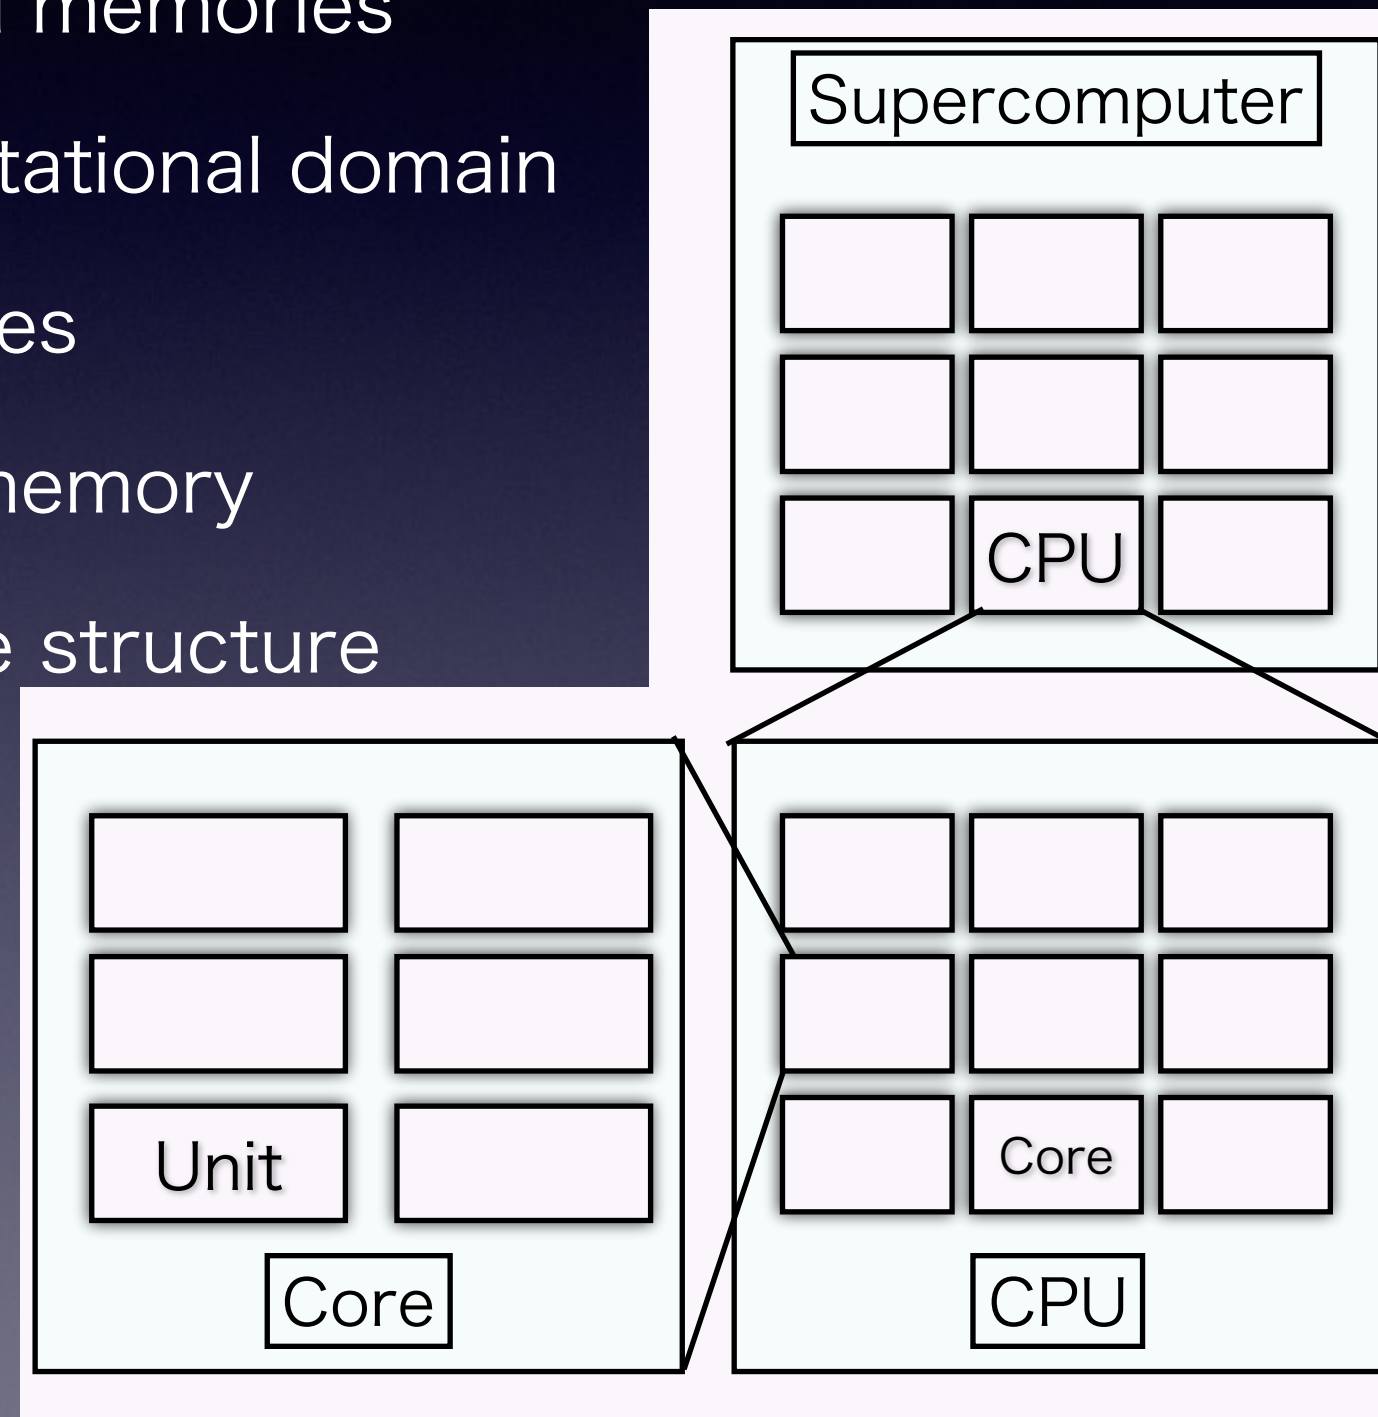
# Why is massively parallel particle simulation needed ?

- Simulations with many particles and long time are desired.
- CPU clock speed will not grow any more.



# Difficulties of massively parallel particle simulations

- Parallelization on distributed memories
  - Decomposition of computational domain
  - Communication of particles
- Parallelization on a shared memory
  - Multi-traverse across tree structure
  - Load balancing
- Parallelization in a core
  - Efficient use of SIMD unit






# Development of tree structure is not easy even for serial processing

- Efficient use of cache memories
- Construction of true structure

$$\frac{d\vec{u}_i}{dt} = \vec{g} \left( \sum_j^N \vec{f}(\vec{u}_i, \vec{u}_j), \vec{u}_i \right)$$

Tree algorithm reduces N to log N in the case of N-body simulation.



# Particle simulation on FDPS

## FDPS

- Domain decomposition
- Exchange of particle data
- Collection of particle data for interaction

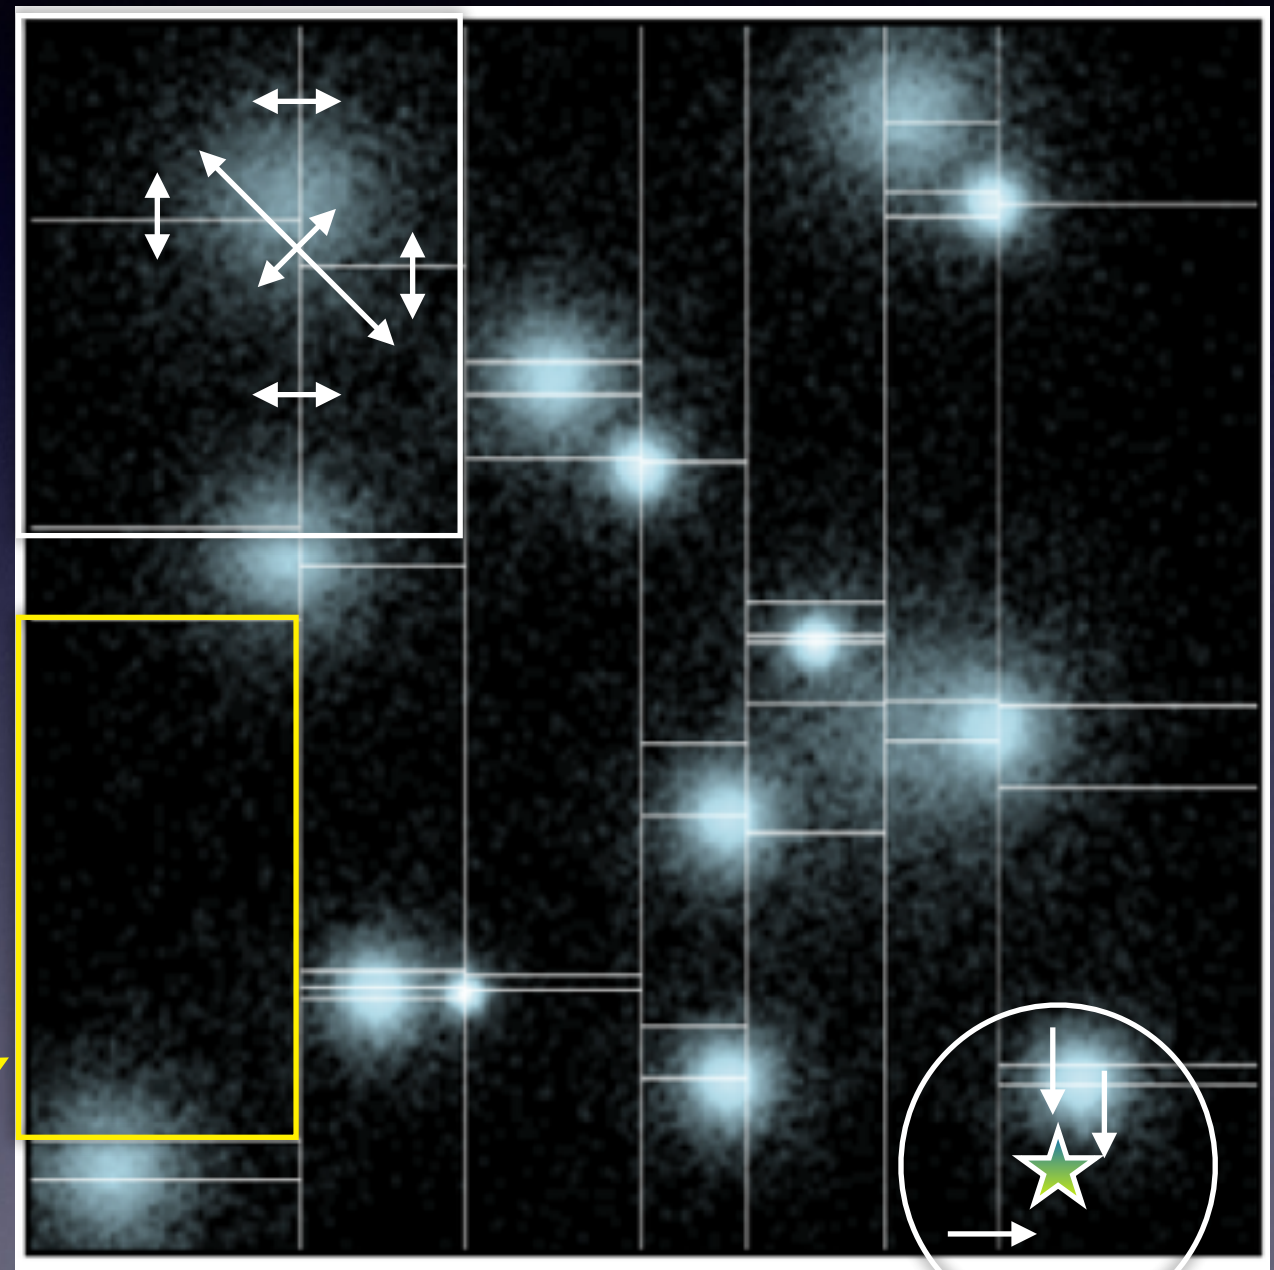
- Actual calculation of interaction
- Integration of particle orbits

## Users

1 MPI process

Domain  
decomposition

Collect  
particle data

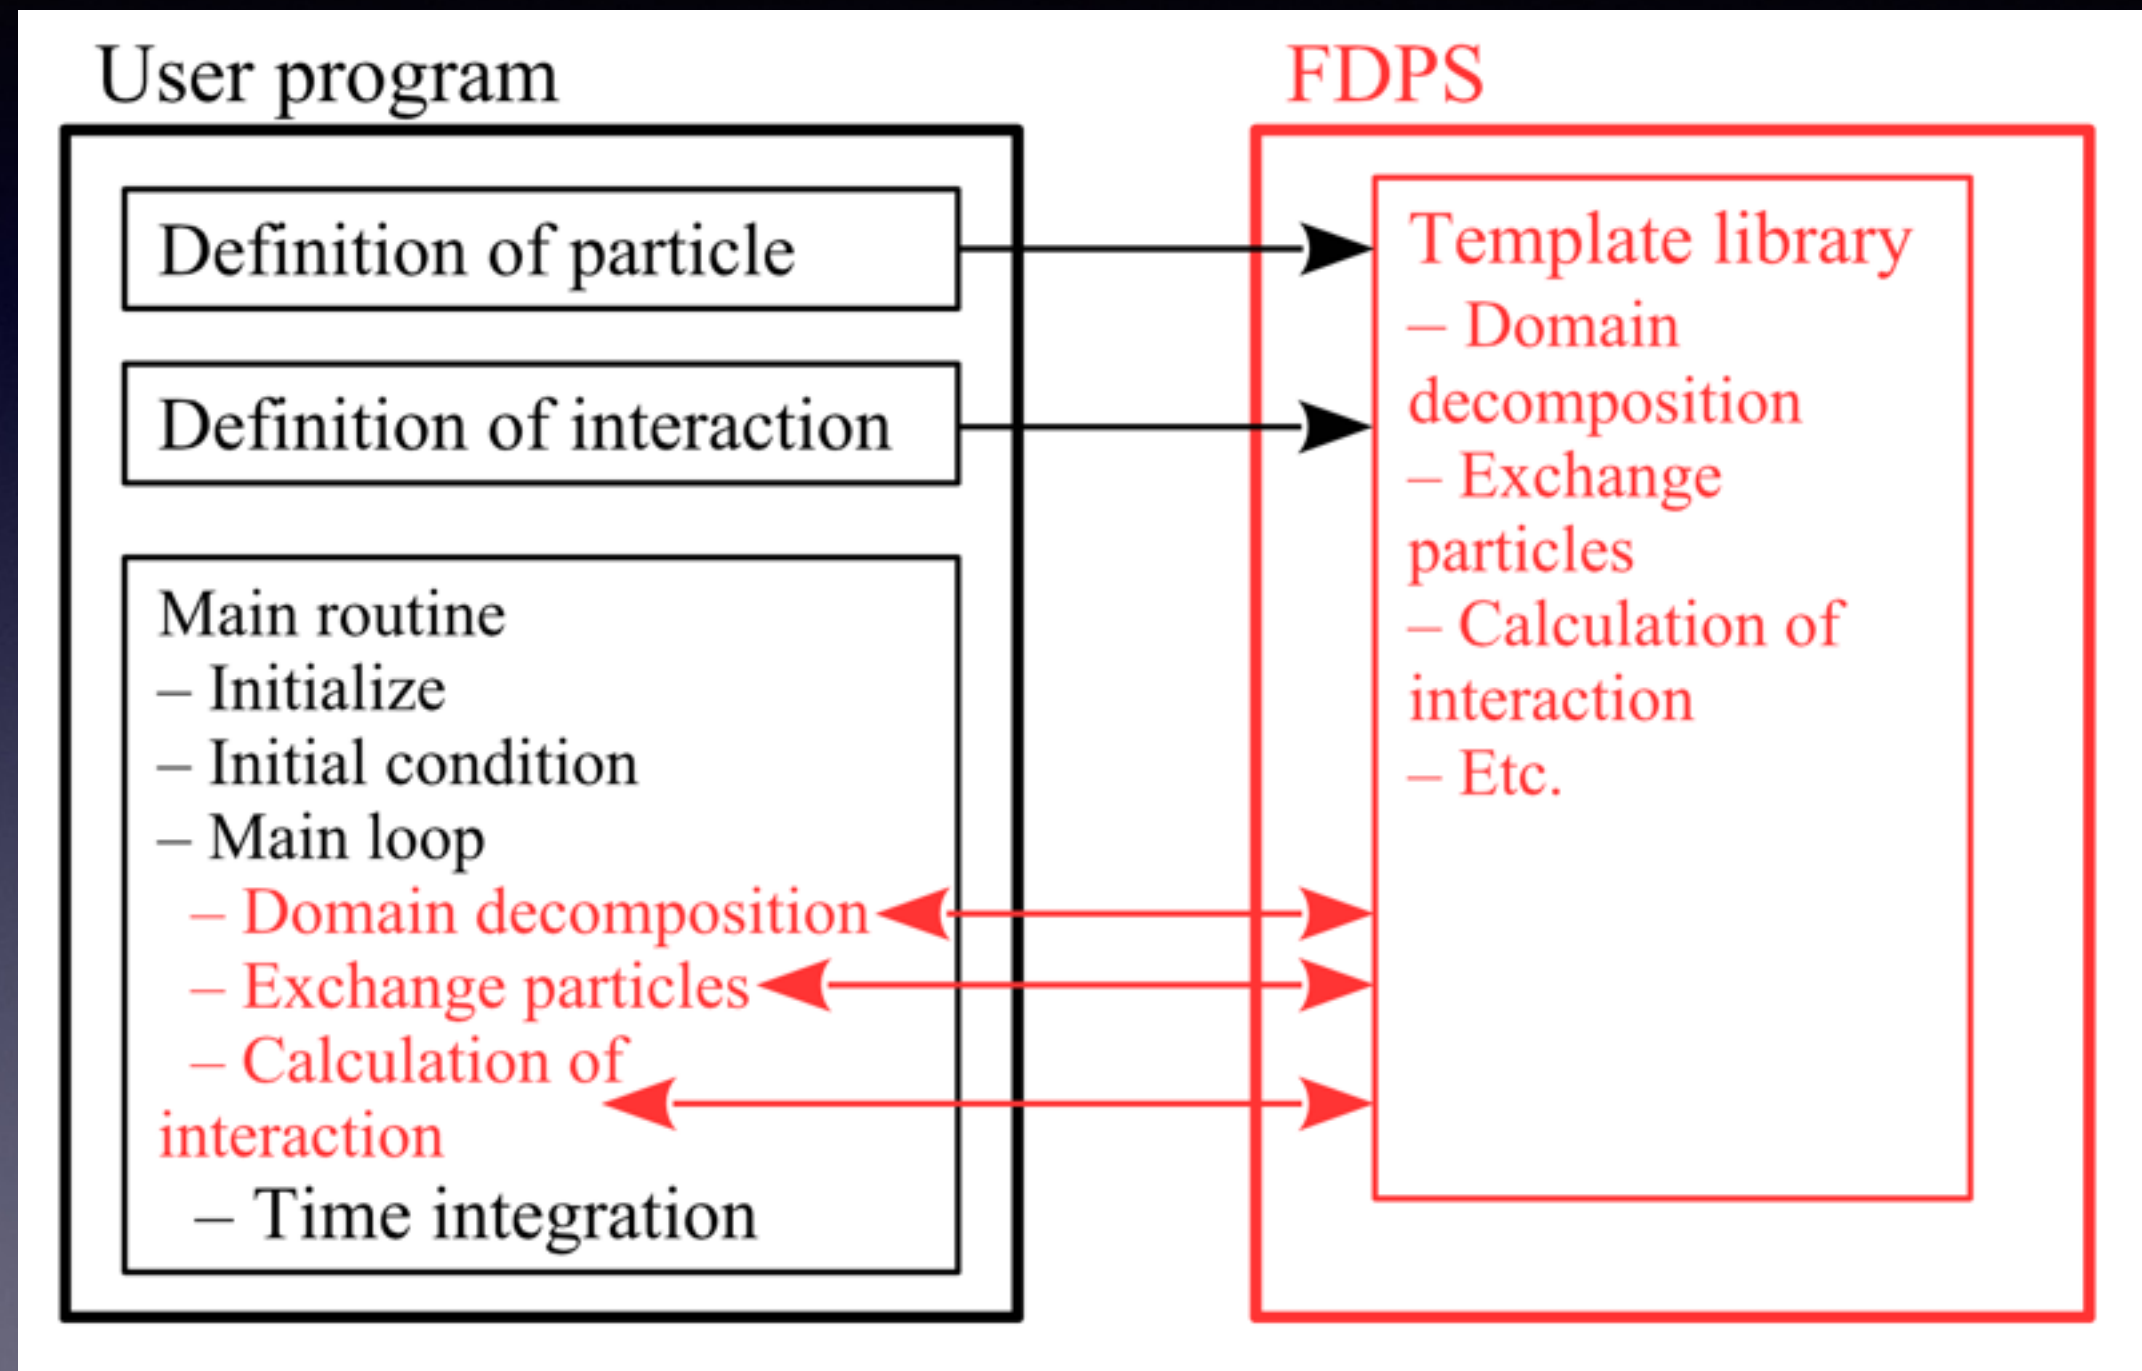


# Overview of FDPS code

- C++ language
  - Class (or structure) for **particle data defined by users**
  - Function pointer (or function object) for **interaction defined by users**
  - Template class to **receive the above class and function pointer**



# Overview of FDPS code





# N-body code

Install FDPS

Definition of particle data

Only 117 lines for  
massively parallel N-  
body code

Users do not need to  
write MPI and OpenMP.

Definition of interaction

Listing 1 shows the complete code which can be actually compiled and run, not only on a single-core machine but also massively-parallel, distributed-memory machines such as the full-node configuration of the K computer. The total number of lines is only 117.

Listing 1: A sample code of N-body simulation

```
1 #include <particle_simulator.hpp>
2 using namespace PS;

3
4 class Nbody{
5 public:
6     F64    mass, eps;
7     F64vec pos, vel, acc;
8     F64vec getPos() const {return pos;}
9     F64 getCharge() const {return mass;}
10    void copyFromFP(const Nbody &in){
11        mass = in.mass;
12        pos = in.pos;
13        eps = in.eps;
14    }
15    void copyFromForce(const Nbody &out) {
16        acc = out.acc;
17    }
18    void clear() {
19        acc = 0.0;
20    }
21    void readAscii(FILE *fp) {
22        fscanf(fp,
23            "%lf%lf%lf%lf%lf%lf%lf%lf",
24            &mass, &eps,
25            &pos.x, &pos.y, &pos.z,
26            &vel.x, &vel.y, &vel.z);
27    }
28    void predict(F64 dt) {
29        vel += (0.5 * dt) * acc;
30        pos += dt * vel;
31    }
32    void correct(F64 dt) {
33        vel += (0.5 * dt) * acc;
34    }
35 };

36
37 template <class TPJ>
38 struct CalcGrav{
39     void operator () (const Nbody * ip,
40                     const S32 ni,
41                     const TPJ * jp,
42                     const S32 nj,
43                     Nbody * force) {
44         for(S32 i=0; i<ni; i++){
45             F64vec xi = ip[i].pos;
46             F64    ep2 = ip[i].eps
47                 * ip[i].eps;
48             F64vec ai = 0.0;
49             for(S32 j=0; j<nj; j++){
50                 F64vec xj = jp[j].pos;
51                 F64vec dr = xi - xj;
52                 F64    mj = jp[j].mass;
53                 F64    dr2 = dr * dr + ep2;
54                 F64    dri = 1.0 / sqrt(dr2);
55                 ai -= (dri * dri * dri
```

```
16         * mj) * dr;
17     }
18     force[i].acc += ai;
19 }
20 };
21
22 template<class Tpsys>
23 void predict(Tpsys &p,
24             const F64 dt) {
25     S32 n = p.getNumberOfParticleLocal();
26     for(S32 i = 0; i < n; i++)
27         p[i].predict(dt);
28 }
29
30 template<class Tpsys>
31 void correct(Tpsys &p,
32             const F64 dt) {
33     S32 n = p.getNumberOfParticleLocal();
34     for(S32 i = 0; i < n; i++)
35         p[i].correct(dt);
36 }
37
38 template <class TDI, class TPS, class TTFP>
39 void calcGravAllAndWriteBack(TDI &dinfo,
40                             TPS &ptcl,
41                             TTFP &tree) {
42     dinfo.decomposeDomainAll(ptcl);
43     ptcl.exchangeParticle(dinfo);
44     tree.calcForceAllAndWriteBack
45         (CalcGrav<Nbody>(),
46         CalcGrav<SPJMonopole>(),
47         ptcl, dinfo);
48 }
49
50 int main(int argc, char *argv[]) {
51     F32 time = 0.0;
52     const F32 tend = 10.0;
53     const F32 dt = 1.0 / 128.0;
54     PS::Initialize(argc, argv);
55     PS::DomainInfo dinfo;
56     dinfo.initialize();
57     PS::ParticleSystem<Nbody> ptcl;
58     ptcl.initialize();
59     PS::TreeForForceLong<Nbody, Nbody,
60         Nbody>::Monopole grav;
61     grav.initialize(0);
62     ptcl.readParticleAscii(argv[1]);
63     calcGravAllAndWriteBack(dinfo,
64                             ptcl,
65                             grav);
66
67     while(time < tend) {
68         predict(ptcl, dt);
69         calcGravAllAndWriteBack(dinfo,
70                                 ptcl,
71                                 grav);
72
73         correct(ptcl, dt);
74         time += dt;
75     }
76     PS::Finalize();
77     return 0;
78 }
```

# Definition of particle data

Particle data:

- mass
- position
- velocity
- acceleration
- gravitational softening

Communication with FDPS:

- Which is position?
- Which is mass?
- Which are required for interaction?
- Which are results?

```
4 class Nbody{
5 public:
6     F64    mass, eps;
7     F64vec pos, vel, acc;
8     F64vec getPos() const {return pos;}
9     F64    getCharge() const {return mass;}
10    void copyFromFP(const Nbody &in){
11        mass = in.mass;
12        pos  = in.pos;
13        eps  = in.eps;
14    }
15    void copyFromForce(const Nbody &out) {
16        acc = out.acc;
17    }
18    void clear() {
19        acc = 0.0;
20    }
21    void readAscii(FILE *fp) {
22        fscanf(fp,
23            "%lf%lf%lf%lf%lf%lf%lf%lf",
24            &mass, &eps,
25            &pos.x, &pos.y, &pos.z,
26            &vel.x, &vel.y, &vel.z);
27    }
28    void predict(F64 dt) {
29        vel += (0.5 * dt) * acc;
30        pos += dt * vel;
31    }
32    void correct(F64 dt) {
33        vel += (0.5 * dt) * acc;
34    }
35 };
```



# Definition of interaction

- Array of i-particle data
- The number of i-particles
- Array of j-particle data
- The number of j-particles
- Array to store the results

i-particle: particle receiving force

j-particle: particle exerting force

Calculate gravitational  
force on i-particle

```
37 template <class TPJ>
38 struct CalcGrav{
39     void operator ()(const Nbody * ip,
40                     const S32 ni,
41                     const TPJ * jp,
42                     const S32 nj,
43                     Nbody * force) {
44         for(S32 i=0; i<ni; i++){
45             F64vec xi = ip[i].pos;
46             F64      ep2 = ip[i].eps
47                 * ip[i].eps;
48             F64vec ai = 0.0;
49             for(S32 j=0; j<nj; j++){
50                 F64vec xj = jp[j].pos;
51                 F64vec dr = xi - xj;
52                 F64 mj = jp[j].mass;
53                 F64 dr2 = dr * dr + ep2;
54                 F64 dri = 1.0 / sqrt(dr2);
55                 ai -= (dri * dri * dri
56                     * mj) * dr;
57             }
58             force[i].acc += ai;
59         }
60     }
61 };
```

# Main routine

Template class for exchange of particle data

Template class for interaction calculation

```
template <class TDI, class TPS, class TTFF>
void calcGravAllAndWriteBack(TDI &dinfo,
                             TPS &ptcl,
                             TTFF &tree) {
    dinfo.decomposeDomainAll(ptcl);
    ptcl.exchangeParticle(dinfo);
    tree.calcForceAllAndWriteBack
        (CalcGrav<Nbody>(),
         CalcGrav<SPJMonopole>(),
         ptcl, dinfo);
}
```



```
91 int main(int argc, char *argv[]) {
92     F32 time = 0.0;
93     const F32 tend = 10.0;
94     const F32 dtm = 1.0 / 128.0;
95     PS::Initialize(argc, argv);
96     PS::DomainInfo dinfo;
97     dinfo.initialize();
98     PS::ParticleSystem<Nbody> ptcl;
99     ptcl.initialize();
100    PS::TreeForForceLong<Nbody, Nbody>
101        <Nbody>::Monopole grav;
102    grav.initialize(0);
103    ptcl.readParticleAscii(argv[1]);
104    calcGravAllAndWriteBack(dinfo,
105                             ptcl,
106                             grav);
107    while(time < tend) {
108        predict(ptcl, dtm);
109        calcGravAllAndWriteBack(dinfo,
110                                 ptcl,
111                                 grav);
112        correct(ptcl, dtm);
113        time += dtm;
114    }
115    PS::Finalize();
116    return 0;
117 }
```



# N-body on K computer


- Plummer model
- $N: 2.1 \times 10^6/\text{node}$
- Accuracy:  $\Theta=0.4$ , up to quadrupole

Example of N-body

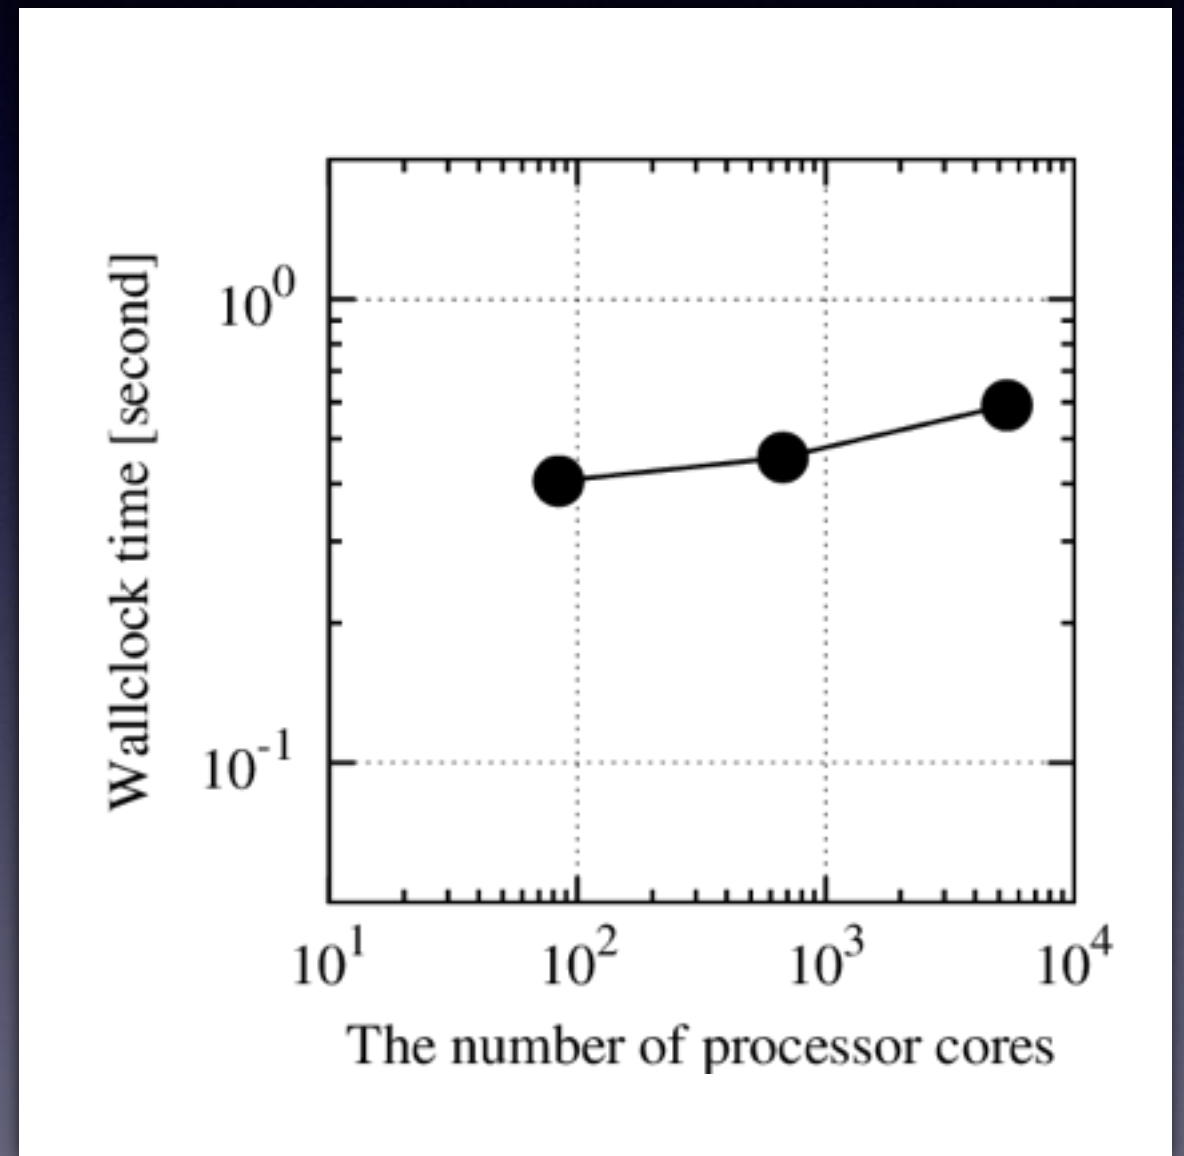


# SPH performance @ CfCA XC30

## (Merging binary white dwarfs)




Tanikawa et al. (2015)




N: 8000/core

# Applications


Spiral galaxy  
(N-body)




Merging  
binary  
white  
dwarfs  
(SPH)



Cosmology  
(N-body with  
periodic  
boundary)



Giant  
impact  
(SPH)





# Sample code

- Gravitational N-body, open boundary
- SPH, ideal EoS
- SPH, ideal EoS, gravity
- Van der Waals
- Gravitational N-body, periodic boundary  
(Yoshikawa)



# N-body on Xeon Phi

- Benchmark
  - N: 1 million
  - Tree accuracy:  $\theta = 0.5$ , up to monopole
  - Interaction accuracy: single precision
  - inverse square root accuracy: 23-24bit
- Computers
  - Xeon Phi 5110P (Performance 1TFlops, TDP 225W)
  - XC30 1 node (Performance 1TFlops, TDP 270W)
- Wallclock time per step
  - Xeon Phi: 0.35 second
  - XC30: 0.21 second
- Caution
  - FDPS has been not yet tuned for many-core processors.

# Publication

- <https://github.com/FDPS/FDPS>
  - サポート: fdps-support@mail.jmlab.jp
- Iwasawa, Tanikawa, Hosono, Nitadori, Muranushi, Makino, WOLFHPC 2015: proceedings of Fifth International Workshop on Domain-Specific Languages and High-Level Frameworks for High Performance Computing
- FDPS講習会 2015/07/22 @ 神戸

# Future work

- FDPS 2.0
  - Support APIs to calculate interactions on many-core processors, such as GPU and PEZY-SC.
- FDPS ?.0
  - All on many-core processors.

# Summary

- <https://github.com/FDPS/FDPS>
- We have developed FDPS.
- You can develop particle simulation codes easily.
- N-body code can be developed only in 117 lines.
- N-body performance is 50% of theoretical peak performance of K computer.
- Weak scaling of SPH is good on CfCA XC30.